

C4W4_Assignment

April 5, 2023

1 Week 4: Using real world data

Welcome! So far you have worked exclusively with generated data. This time you will be using the [Daily Minimum Temperatures in Melbourne](#) dataset which contains data of the daily minimum temperatures recorded in Melbourne from 1981 to 1990. In addition to be using Tensorflow's layers for processing sequence data such as Recurrent layers or LSTMs you will also use Convolutional layers to improve the model's performance.

Let's get started!

```
[1]: import csv
import pickle
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from dataclasses import dataclass
```

Begin by looking at the structure of the csv that contains the data:

```
[2]: TEMPERATURES_CSV = './data/daily-min-temperatures.csv'

with open(TEMPERATURES_CSV, 'r') as csvfile:
    print(f"Header looks like this:\n\n{csvfile.readline()}")
    print(f"First data point looks like this:\n\n{csvfile.readline()}")
    print(f"Second data point looks like this:\n\n{csvfile.readline()}")
```

Header looks like this:

```
"Date","Temp"
```

First data point looks like this:

```
"1981-01-01",20.7
```

Second data point looks like this:

```
"1981-01-02",17.9
```

As you can see, each data point is composed of the date and the recorded minimum temperature for that date.

In the first exercise you will code a function to read the data from the csv but for now run the next cell to load a helper function to plot the time series.

```
[3]: def plot_series(time, series, format="-", start=0, end=None):  
    plt.plot(time[start:end], series[start:end], format)  
    plt.xlabel("Time")  
    plt.ylabel("Value")  
    plt.grid(True)
```

1.1 Parsing the raw data

Now you need to read the data from the csv file. To do so, complete the `parse_data_from_file` function.

A couple of things to note:

- You should omit the first line as the file contains headers.
- There is no need to save the data points as numpy arrays, regular lists is fine.
- To read from csv files use `csv.reader` by passing the appropriate arguments.
- `csv.reader` returns an iterable that returns each row in every iteration. So the temperature can be accessed via `row[1]` and the date can be discarded.
- The `times` list should contain every timestep (starting at zero), which is just a sequence of ordered numbers with the same length as the `temperatures` list.
- The values of the `temperatures` should be of `float` type. You can use Python's built-in `float` function to ensure this.

```
[6]: # From mysef  
import pandas as pd  
df = pd.read_csv("../data/daily-min-temperatures.csv")  
df.head()
```

```
[6]:      Date  Temp  
0  1981-01-01  20.7  
1  1981-01-02  17.9  
2  1981-01-03  18.8  
3  1981-01-04  14.6  
4  1981-01-05  15.8
```

```
[23]: def parse_data_from_file(filename):  
  
    times = []  
    temperatures = []  
  
    with open(filename) as csvfile:  
  
        ### START CODE HERE
```

```

        reader = csv.reader(csvfile, delimiter=',')
        next(reader)

        for i, row in enumerate(reader):
#            times.append(int(row[0]))
            times.append(i)
            temperatures.append(float(row[1]))

        times = np.array(times)
        temperatures = np.array(temperatures)

        ### END CODE HERE

    return times, temperatures

```


The next cell will use your function to compute the `times` and `temperatures` and will save these as numpy arrays within the `G` dataclass. This cell will also plot the time series:

```

[24]: # Test your function and save all "global" variables within the G class (G_
      ↪stands for global)
      @dataclass
      class G:
          TEMPERATURES_CSV = './data/daily-min-temperatures.csv'
          times, temperatures = parse_data_from_file(TEMPERATURES_CSV)
          TIME = np.array(times)
          SERIES = np.array(temperatures)
          SPLIT_TIME = 2500
          WINDOW_SIZE = 64
          BATCH_SIZE = 32
          SHUFFLE_BUFFER_SIZE = 1000

      plt.figure(figsize=(10, 6))
      plot_series(G.TIME, G.SERIES)
      plt.show()

```



Expected Output:

1.2 Processing the data

Since you already coded the `train_val_split` and `windowed_dataset` functions during past week's assignments, this time they are provided for you:

```
[25]: def train_val_split(time, series, time_step=G.SPLIT_TIME):

    time_train = time[:time_step]
    series_train = series[:time_step]
    time_valid = time[time_step:]
    series_valid = series[time_step:]

    return time_train, series_train, time_valid, series_valid

# Split the dataset
time_train, series_train, time_valid, series_valid = train_val_split(G.TIME, G.
↪SERIES)
```

```
[26]: def windowed_dataset(series, window_size=G.WINDOW_SIZE, batch_size=G.
↪BATCH_SIZE, shuffle_buffer=G.SHUFFLE_BUFFER_SIZE):
    ds = tf.data.Dataset.from_tensor_slices(series)
    ds = ds.window(window_size + 1, shift=1, drop_remainder=True)
```

```

ds = ds.flat_map(lambda w: w.batch(window_size + 1))
ds = ds.shuffle(shuffle_buffer)
ds = ds.map(lambda w: (w[:-1], w[-1]))
ds = ds.batch(batch_size).prefetch(1)
return ds

```

Apply the transformation to the training set

```

train_set = windowed_dataset(series_train, window_size=G.WINDOW_SIZE,
    ↪ batch_size=G.BATCH_SIZE, shuffle_buffer=G.SHUFFLE_BUFFER_SIZE)

```

Metal device set to: Apple M1 Pro

```

2023-04-05 17:00:48.901486: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2023-04-05 17:00:48.901877: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)

```

1.3 Defining the model architecture

Now that you have a function that will process the data before it is fed into your neural network for training, it is time to define your layer architecture. Just as in last week's assignment you will do the layer definition and compilation in two separate steps. Begin by completing the `create_uncompiled_model` function below.

This is done so you can reuse your model's layers for the learning rate adjusting and the actual training.

Hint:

- Lambda layers are not required.
- Use a combination of Conv1D and LSTM layers followed by Dense layers

```

[50]: def create_uncompiled_model():

    ### START CODE HERE

    model = tf.keras.models.Sequential([
        tf.keras.layers.Conv1D(filters=60, kernel_size=5,
                                strides=1,
                                activation="relu",
                                padding='causal',
                                input_shape=[G.WINDOW_SIZE, 1]),
        tf.keras.layers.LSTM(64, return_sequences=True),
        tf.keras.layers.LSTM(64),

```

```

        tf.keras.layers.Dense(30, activation="relu"),
        tf.keras.layers.Dense(10, activation="relu"),
        tf.keras.layers.Dense(1),
        tf.keras.layers.Lambda(lambda x: x * 400)
    ])

    ### END CODE HERE

    return model

```

```

[51]: # Test your uncompiled model
uncompiled_model = create_uncompiled_model()

try:
    uncompiled_model.predict(train_set)
except:
    print("Your current architecture is incompatible with the windowed dataset,
    ↪ try adjusting it.")
else:
    print("Your current architecture is compatible with the windowed dataset! :
    ↪ ")

```

```

2023-04-05 17:19:16.215099: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:19:16.337436: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

7/Unknown - 1s 18ms/step

2023-04-05 17:19:16.508583: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

77/77 [=====] - 2s 20ms/step
Your current architecture is compatible with the windowed dataset! :)

```

1.4 Adjusting the learning rate - (Optional Exercise)

As you saw in the lecture you can leverage Tensorflow's callbacks to dynamically vary the learning rate during training. This can be helpful to get a better sense of which learning rate better accommodates to the problem at hand.

Notice that this is only changing the learning rate during the training process to give you an idea of what a reasonable learning rate is and should not be confused with selecting the best learning rate, this is known as hyperparameter optimization and it is outside the scope of this course.

For the optimizers you can try out:

- tf.keras.optimizers.Adam
- tf.keras.optimizers.SGD with a momentum of 0.9

```
[47]: def adjust_learning_rate(dataset):

    model = create_uncompiled_model()

    lr_schedule = tf.keras.callbacks.LearningRateScheduler(lambda epoch: 1e-4 *
↳ 10**(epoch / 20))

    ### START CODE HERE

    # Select your optimizer
    optimizer = tf.keras.optimizers.SGD(momentum=0.9)

    # Compile the model passing in the appropriate loss
    model.compile(loss=tf.keras.losses.Huber(),
                  optimizer=optimizer,
                  metrics=["mae"])

    ### END CODE HERE

    history = model.fit(dataset, epochs=100, callbacks=[lr_schedule])

    return history
```

```
[48]: # Run the training with dynamic LR
lr_history = adjust_learning_rate(train_set)
```

Epoch 1/100

```
2023-04-05 17:14:52.759805: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:14:53.010419: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:14:53.105323: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:14:53.235257: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:14:53.396429: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

77/77 [=====] - 4s 38ms/step - loss: 13.6027 - mae:
14.0893 - lr: 1.0000e-04
```

Epoch 2/100
77/77 [=====] - 3s 32ms/step - loss: 6.0536 - mae: 6.5378 - lr: 1.1220e-04

Epoch 3/100
77/77 [=====] - 3s 32ms/step - loss: 14.2872 - mae: 14.7816 - lr: 1.2589e-04

Epoch 4/100
77/77 [=====] - 3s 32ms/step - loss: 20.9056 - mae: 21.4046 - lr: 1.4125e-04

Epoch 5/100
77/77 [=====] - 3s 33ms/step - loss: 25.9744 - mae: 26.4730 - lr: 1.5849e-04

Epoch 6/100
77/77 [=====] - 3s 33ms/step - loss: 24.8408 - mae: 25.3397 - lr: 1.7783e-04

Epoch 7/100
77/77 [=====] - 3s 33ms/step - loss: 20.3105 - mae: 20.8050 - lr: 1.9953e-04

Epoch 8/100
77/77 [=====] - 3s 32ms/step - loss: 22.8652 - mae: 23.3622 - lr: 2.2387e-04

Epoch 9/100
77/77 [=====] - 3s 32ms/step - loss: 20.7610 - mae: 21.2597 - lr: 2.5119e-04

Epoch 10/100
77/77 [=====] - 3s 33ms/step - loss: 23.1088 - mae: 23.6088 - lr: 2.8184e-04

Epoch 11/100
77/77 [=====] - 3s 32ms/step - loss: 33.2279 - mae: 33.7269 - lr: 3.1623e-04

Epoch 12/100
77/77 [=====] - 3s 32ms/step - loss: 33.8599 - mae: 34.3587 - lr: 3.5481e-04

Epoch 13/100
77/77 [=====] - 2s 31ms/step - loss: 63.8486 - mae: 64.3458 - lr: 3.9811e-04

Epoch 14/100
77/77 [=====] - 3s 32ms/step - loss: 108.0357 - mae: 108.5356 - lr: 4.4668e-04

Epoch 15/100
77/77 [=====] - 3s 33ms/step - loss: 112.2182 - mae: 112.7179 - lr: 5.0119e-04

Epoch 16/100
77/77 [=====] - 3s 32ms/step - loss: 72.2861 - mae: 72.7846 - lr: 5.6234e-04

Epoch 17/100
77/77 [=====] - 3s 32ms/step - loss: 45.9963 - mae: 46.4951 - lr: 6.3096e-04

Epoch 18/100
77/77 [=====] - 3s 32ms/step - loss: 100.5617 - mae: 101.0601 - lr: 7.0795e-04

Epoch 19/100
77/77 [=====] - 3s 32ms/step - loss: 115.9424 - mae: 116.4412 - lr: 7.9433e-04

Epoch 20/100
77/77 [=====] - 3s 33ms/step - loss: 162.5921 - mae: 163.0902 - lr: 8.9125e-04

Epoch 21/100
77/77 [=====] - 3s 33ms/step - loss: 116.1965 - mae: 116.6952 - lr: 0.0010

Epoch 22/100
77/77 [=====] - 3s 32ms/step - loss: 92.6116 - mae: 93.1116 - lr: 0.0011

Epoch 23/100
77/77 [=====] - 3s 33ms/step - loss: 105.4148 - mae: 105.9148 - lr: 0.0013

Epoch 24/100
77/77 [=====] - 3s 33ms/step - loss: 243.5321 - mae: 244.0316 - lr: 0.0014

Epoch 25/100
77/77 [=====] - 3s 32ms/step - loss: 273.9176 - mae: 274.4176 - lr: 0.0016

Epoch 26/100
77/77 [=====] - 3s 32ms/step - loss: 270.6005 - mae: 271.1005 - lr: 0.0018

Epoch 27/100
77/77 [=====] - 3s 33ms/step - loss: 300.8000 - mae: 301.2999 - lr: 0.0020

Epoch 28/100
77/77 [=====] - 3s 32ms/step - loss: 345.3090 - mae: 345.8090 - lr: 0.0022

Epoch 29/100
77/77 [=====] - 3s 32ms/step - loss: 301.0367 - mae: 301.5365 - lr: 0.0025

Epoch 30/100
77/77 [=====] - 3s 33ms/step - loss: 366.7590 - mae: 367.2589 - lr: 0.0028

Epoch 31/100
77/77 [=====] - 3s 32ms/step - loss: 545.8618 - mae: 546.3618 - lr: 0.0032

Epoch 32/100
77/77 [=====] - 3s 32ms/step - loss: 803.4770 - mae: 803.9770 - lr: 0.0035

Epoch 33/100
77/77 [=====] - 3s 33ms/step - loss: 437.2830 - mae: 437.7826 - lr: 0.0040

Epoch 34/100
77/77 [=====] - 3s 33ms/step - loss: 555.5274 - mae: 556.0274 - lr: 0.0045

Epoch 35/100
77/77 [=====] - 3s 33ms/step - loss: 908.8295 - mae: 909.3295 - lr: 0.0050

Epoch 36/100
77/77 [=====] - 3s 33ms/step - loss: 469.8769 - mae: 470.3767 - lr: 0.0056

Epoch 37/100
77/77 [=====] - 3s 33ms/step - loss: 789.3442 - mae: 789.8442 - lr: 0.0063

Epoch 38/100
77/77 [=====] - 3s 33ms/step - loss: 585.9149 - mae: 586.4147 - lr: 0.0071

Epoch 39/100
77/77 [=====] - 3s 33ms/step - loss: 1159.3988 - mae: 1159.8988 - lr: 0.0079

Epoch 40/100
77/77 [=====] - 3s 33ms/step - loss: 2455.3201 - mae: 2455.8201 - lr: 0.0089

Epoch 41/100
77/77 [=====] - 3s 33ms/step - loss: 1222.2289 - mae: 1222.7289 - lr: 0.0100

Epoch 42/100
77/77 [=====] - 3s 33ms/step - loss: 1297.4390 - mae: 1297.9390 - lr: 0.0112

Epoch 43/100
77/77 [=====] - 3s 33ms/step - loss: 1508.9277 - mae: 1509.4277 - lr: 0.0126

Epoch 44/100
77/77 [=====] - 3s 34ms/step - loss: 1805.8767 - mae: 1806.3767 - lr: 0.0141

Epoch 45/100
77/77 [=====] - 3s 34ms/step - loss: 2444.3838 - mae: 2444.8838 - lr: 0.0158

Epoch 46/100
77/77 [=====] - 3s 33ms/step - loss: 2169.3372 - mae: 2169.8372 - lr: 0.0178

Epoch 47/100
77/77 [=====] - 3s 33ms/step - loss: 3053.3896 - mae: 3053.8896 - lr: 0.0200

Epoch 48/100
77/77 [=====] - 3s 33ms/step - loss: 3011.2454 - mae: 3011.7454 - lr: 0.0224

Epoch 49/100
77/77 [=====] - 3s 34ms/step - loss: 4328.3726 - mae: 4328.8726 - lr: 0.0251

Epoch 50/100
77/77 [=====] - 3s 33ms/step - loss: 3597.4106 - mae: 3597.9106 - lr: 0.0282

Epoch 51/100
77/77 [=====] - 3s 33ms/step - loss: 5495.2412 - mae: 5495.7412 - lr: 0.0316

Epoch 52/100
77/77 [=====] - 3s 33ms/step - loss: 4519.6250 - mae: 4520.1250 - lr: 0.0355

Epoch 53/100
77/77 [=====] - 3s 33ms/step - loss: 4812.1050 - mae: 4812.6050 - lr: 0.0398

Epoch 54/100
77/77 [=====] - 3s 33ms/step - loss: 12095.2236 - mae: 12095.7236 - lr: 0.0447

Epoch 55/100
77/77 [=====] - 3s 33ms/step - loss: 6053.4116 - mae: 6053.9116 - lr: 0.0501

Epoch 56/100
77/77 [=====] - 3s 33ms/step - loss: 7130.0830 - mae: 7130.5830 - lr: 0.0562

Epoch 57/100
77/77 [=====] - 3s 33ms/step - loss: 5180.4404 - mae: 5180.9404 - lr: 0.0631

Epoch 58/100
77/77 [=====] - 3s 33ms/step - loss: 5937.8276 - mae: 5938.3276 - lr: 0.0708

Epoch 59/100
77/77 [=====] - 3s 34ms/step - loss: 6622.0117 - mae: 6622.5117 - lr: 0.0794

Epoch 60/100
77/77 [=====] - 3s 33ms/step - loss: 7433.9951 - mae: 7434.4951 - lr: 0.0891

Epoch 61/100
77/77 [=====] - 3s 33ms/step - loss: 7811.1865 - mae: 7811.6865 - lr: 0.1000

Epoch 62/100
77/77 [=====] - 3s 33ms/step - loss: 15678.5469 - mae: 15679.0459 - lr: 0.1122

Epoch 63/100
77/77 [=====] - 3s 33ms/step - loss: 31433.2637 - mae: 31433.7676 - lr: 0.1259

Epoch 64/100
77/77 [=====] - 3s 33ms/step - loss: 35283.0781 - mae: 35283.5781 - lr: 0.1413

Epoch 65/100
77/77 [=====] - 3s 33ms/step - loss: 19038.4961 - mae: 19038.9961 - lr: 0.1585

Epoch 66/100
77/77 [=====] - 3s 33ms/step - loss: 32036.6602 - mae: 32037.1602 - lr: 0.1778

Epoch 67/100
77/77 [=====] - 3s 33ms/step - loss: 25470.2715 - mae: 25470.7715 - lr: 0.1995

Epoch 68/100
77/77 [=====] - 3s 33ms/step - loss: 27861.3262 - mae: 27861.8262 - lr: 0.2239

Epoch 69/100
77/77 [=====] - 3s 33ms/step - loss: 45572.5742 - mae: 45573.0781 - lr: 0.2512

Epoch 70/100
77/77 [=====] - 3s 34ms/step - loss: 53643.3242 - mae: 53643.8242 - lr: 0.2818

Epoch 71/100
77/77 [=====] - 3s 34ms/step - loss: 78977.4297 - mae: 78977.9297 - lr: 0.3162

Epoch 72/100
77/77 [=====] - 3s 33ms/step - loss: 86605.7109 - mae: 86606.2109 - lr: 0.3548

Epoch 73/100
77/77 [=====] - 3s 33ms/step - loss: 58110.9023 - mae: 58111.4102 - lr: 0.3981

Epoch 74/100
77/77 [=====] - 3s 34ms/step - loss: 76405.3047 - mae: 76405.7969 - lr: 0.4467

Epoch 75/100
77/77 [=====] - 3s 34ms/step - loss: 73489.0078 - mae: 73489.5078 - lr: 0.5012

Epoch 76/100
77/77 [=====] - 3s 34ms/step - loss: 93734.9062 - mae: 93735.4062 - lr: 0.5623

Epoch 77/100
77/77 [=====] - 3s 33ms/step - loss: 157222.2500 - mae: 157222.7031 - lr: 0.6310

Epoch 78/100
77/77 [=====] - 3s 33ms/step - loss: 131304.7500 - mae: 131305.2344 - lr: 0.7079

Epoch 79/100
77/77 [=====] - 3s 34ms/step - loss: 221897.7031 - mae: 221898.2188 - lr: 0.7943

Epoch 80/100
77/77 [=====] - 3s 34ms/step - loss: 226214.2500 - mae: 226214.7344 - lr: 0.8913

Epoch 81/100
77/77 [=====] - 3s 34ms/step - loss: 185135.6250 - mae: 185136.1250 - lr: 1.0000

Epoch 82/100
77/77 [=====] - 3s 33ms/step - loss: 177837.0000 - mae: 177837.5625 - lr: 1.1220

Epoch 83/100
77/77 [=====] - 3s 33ms/step - loss: 259452.5312 - mae: 259453.1719 - lr: 1.2589

Epoch 84/100
77/77 [=====] - 3s 34ms/step - loss: 189023.5312 - mae: 189024.0469 - lr: 1.4125

Epoch 85/100
77/77 [=====] - 3s 33ms/step - loss: 66878.5859 - mae: 66879.0859 - lr: 1.5849

Epoch 86/100
77/77 [=====] - 3s 35ms/step - loss: 264427.2188 - mae: 264427.6875 - lr: 1.7783

Epoch 87/100
77/77 [=====] - 3s 34ms/step - loss: 460274.7500 - mae: 460275.1875 - lr: 1.9953

Epoch 88/100
77/77 [=====] - 3s 33ms/step - loss: 826394.1250 - mae: 826394.5000 - lr: 2.2387

Epoch 89/100
77/77 [=====] - 3s 33ms/step - loss: 709319.6875 - mae: 709320.1875 - lr: 2.5119

Epoch 90/100
77/77 [=====] - 3s 33ms/step - loss: 718335.0625 - mae: 718335.4375 - lr: 2.8184

Epoch 91/100
77/77 [=====] - 3s 33ms/step - loss: 522067.7500 - mae: 522068.3438 - lr: 3.1623

Epoch 92/100
77/77 [=====] - 3s 34ms/step - loss: 624183.0625 - mae: 624183.6250 - lr: 3.5481

Epoch 93/100
77/77 [=====] - 3s 34ms/step - loss: 489455.6562 - mae: 489456.1250 - lr: 3.9811

Epoch 94/100
77/77 [=====] - 3s 33ms/step - loss: 571648.5000 - mae: 571649.0000 - lr: 4.4668

Epoch 95/100
77/77 [=====] - 3s 33ms/step - loss: 756480.1875 - mae: 756480.8750 - lr: 5.0119


Epoch 96/100
77/77 [=====] - 3s 33ms/step - loss: 674925.6250 - mae: 674926.3125 - lr: 5.6234

Epoch 97/100
77/77 [=====] - 3s 33ms/step - loss: 969099.7500 - mae: 969100.0625 - lr: 6.3096

```
Epoch 98/100
77/77 [=====] - 3s 33ms/step - loss: 935954.9375 - mae:
935955.6875 - lr: 7.0795
Epoch 99/100
77/77 [=====] - 3s 35ms/step - loss: 1210791.7500 -
mae: 1210792.0000 - lr: 7.9433
Epoch 100/100
77/77 [=====] - 3s 34ms/step - loss: 1127579.6250 -
mae: 1127580.0000 - lr: 8.9125
```

```
[49]: plt.semilogx(lr_history.history["lr"], lr_history.history["loss"])
plt.axis([1e-4, 10, 0, 10])
```

```
[49]: (0.0001, 10.0, 0.0, 10.0)
```



1.5 Compiling the model

Now that you have trained the model while varying the learning rate, it is time to do the actual training that will be used to forecast the time series. For this complete the `create_model` function below.

Notice that you are reusing the architecture you defined in the `create_uncompiled_model` earlier. Now you only need to compile this model using the appropriate loss, optimizer (and learning rate).

Hints:

- The training should be really quick so if you notice that each epoch is taking more than a few seconds, consider trying a different architecture.
- If after the first epoch you get an output like this: loss: nan - mae: nan it is very likely that your network is suffering from exploding gradients. This is a common problem if you used SGD as optimizer and set a learning rate that is too high. If you encounter this problem consider lowering the learning rate or using Adam with the default learning rate.

```
[52]: def create_model():

    learning_rate = 1e-5
    model = create_uncompiled_model()

    ### START CODE HERE

    model.compile(loss=tf.keras.losses.Huber(),
                  optimizer=tf.keras.optimizers.
↳SGD(learning_rate=learning_rate, momentum=0.9),
                  metrics=["mae"])

    ### END CODE HERE

    return model
```

```
[53]: # Save an instance of the model
model = create_model()

# Train it
history = model.fit(train_set, epochs=50)
```

Epoch 1/50

```
2023-04-05 17:19:35.366218: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:19:35.616472: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:19:35.706413: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:19:35.834767: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:19:36.008116: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
```

Plugin optimizer for device_type GPU is enabled.

```
77/77 [=====] - 4s 38ms/step - loss: 4.7142 - mae: 5.1861
Epoch 2/50
77/77 [=====] - 3s 32ms/step - loss: 2.2378 - mae: 2.6997
Epoch 3/50
77/77 [=====] - 3s 33ms/step - loss: 2.2571 - mae: 2.7200
Epoch 4/50
77/77 [=====] - 3s 32ms/step - loss: 2.3375 - mae: 2.8013
Epoch 5/50
77/77 [=====] - 3s 33ms/step - loss: 1.8572 - mae: 2.3094
Epoch 6/50
77/77 [=====] - 3s 32ms/step - loss: 1.9440 - mae: 2.3967
Epoch 7/50
77/77 [=====] - 3s 32ms/step - loss: 1.9415 - mae: 2.3961
Epoch 8/50
77/77 [=====] - 3s 32ms/step - loss: 1.8128 - mae: 2.2672
Epoch 9/50
77/77 [=====] - 3s 32ms/step - loss: 2.0640 - mae: 2.5180
Epoch 10/50
77/77 [=====] - 3s 32ms/step - loss: 1.9165 - mae: 2.3720
Epoch 11/50
77/77 [=====] - 3s 32ms/step - loss: 1.8967 - mae: 2.3508
Epoch 12/50
77/77 [=====] - 3s 33ms/step - loss: 1.8364 - mae: 2.2919
Epoch 13/50
77/77 [=====] - 3s 32ms/step - loss: 1.8570 - mae: 2.3133
Epoch 14/50
77/77 [=====] - 3s 32ms/step - loss: 2.0748 - mae: 2.5348
Epoch 15/50
77/77 [=====] - 3s 32ms/step - loss: 1.6699 - mae: 2.1177
Epoch 16/50
77/77 [=====] - 3s 33ms/step - loss: 1.6900 - mae:
```


2.1362
 Epoch 17/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7442 - mae:
 2.1942
 Epoch 18/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8343 - mae:
 2.2874
 Epoch 19/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8807 - mae:
 2.3363
 Epoch 20/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8431 - mae:
 2.2943
 Epoch 21/50
 77/77 [=====] - 3s 34ms/step - loss: 1.7236 - mae:
 2.1741
 Epoch 22/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7677 - mae:
 2.2206
 Epoch 23/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7048 - mae:
 2.1576
 Epoch 24/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7017 - mae:
 2.1520
 Epoch 25/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8600 - mae:
 2.3148
 Epoch 26/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7146 - mae:
 2.1644
 Epoch 27/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8255 - mae:
 2.2806
 Epoch 28/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7239 - mae:
 2.1801
 Epoch 29/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6116 - mae:
 2.0577
 Epoch 30/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8755 - mae:
 2.3325
 Epoch 31/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7397 - mae:
 2.1927
 Epoch 32/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6416 - mae:

2.0876
 Epoch 33/50
 77/77 [=====] - 3s 34ms/step - loss: 1.6999 - mae:
 2.1503
 Epoch 34/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6243 - mae:
 2.0723
 Epoch 35/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6369 - mae:
 2.0844
 Epoch 36/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6273 - mae:
 2.0768
 Epoch 37/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7954 - mae:
 2.2452
 Epoch 38/50
 77/77 [=====] - 3s 34ms/step - loss: 1.8058 - mae:
 2.2550
 Epoch 39/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6549 - mae:
 2.1032
 Epoch 40/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7817 - mae:
 2.2360
 Epoch 41/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6935 - mae:
 2.1454
 Epoch 42/50
 77/77 [=====] - 3s 33ms/step - loss: 1.8124 - mae:
 2.2648
 Epoch 43/50
 77/77 [=====] - 3s 33ms/step - loss: 1.6481 - mae:
 2.0986
 Epoch 44/50
 77/77 [=====] - 3s 34ms/step - loss: 1.7360 - mae:
 2.1883
 Epoch 45/50
 77/77 [=====] - 3s 33ms/step - loss: 1.7709 - mae:
 2.2229
 Epoch 46/50
 77/77 [=====] - 3s 34ms/step - loss: 1.9003 - mae:
 2.3560
 Epoch 47/50
 77/77 [=====] - 3s 34ms/step - loss: 1.7020 - mae:
 2.1523
 Epoch 48/50
 77/77 [=====] - 3s 34ms/step - loss: 1.8370 - mae:

```

2.2896
Epoch 49/50
77/77 [=====] - 3s 33ms/step - loss: 1.5972 - mae:
2.0441
Epoch 50/50
77/77 [=====] - 3s 33ms/step - loss: 1.6221 - mae:
2.0694

```

1.6 Evaluating the forecast

Now it is time to evaluate the performance of the forecast. For this you can use the `compute_metrics` function that you coded in a previous assignment:

```

[54]: def compute_metrics(true_series, forecast):

    mse = tf.keras.metrics.mean_squared_error(true_series, forecast).numpy()
    mae = tf.keras.metrics.mean_absolute_error(true_series, forecast).numpy()

    return mse, mae

```

At this point only the model that will perform the forecast is ready but you still need to compute the actual forecast.

1.7 Faster model forecasts

In the previous week you saw a faster approach compared to using a for loop to compute the forecasts for every point in the sequence. Remember that this faster approach uses batches of data.

The code to implement this is provided in the `model_forecast` below. Notice that the code is very similar to the one in the `windowed_dataset` function with the differences that: - The dataset is windowed using `window_size` rather than `window_size + 1` - No shuffle should be used - No need to split the data into features and labels - A model is used to predict batches of the dataset

```

[55]: def model_forecast(model, series, window_size):
    ds = tf.data.Dataset.from_tensor_slices(series)
    ds = ds.window(window_size, shift=1, drop_remainder=True)
    ds = ds.flat_map(lambda w: w.batch(window_size))
    ds = ds.batch(32).prefetch(1)
    forecast = model.predict(ds)
    return forecast

```

Now compute the actual forecast:

Note: Don't modify the cell below.

The grader uses the same slicing to get the forecast so if you change the cell below you risk having issues when submitting your model for grading.

```


[56]: # Compute the forecast for all the series
rnn_forecast = model_forecast(model, G.SERIES, G.WINDOW_SIZE).squeeze()

```

```
# Slice the forecast to get only the predictions for the validation set
rnn_forecast = rnn_forecast[G.SPLIT_TIME - G.WINDOW_SIZE:-1]

# Plot the forecast
plt.figure(figsize=(10, 6))
plot_series(time_valid, series_valid)
plot_series(time_valid, rnn_forecast)
```

```
2023-04-05 17:21:47.083441: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:21:47.156335: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
2023-04-05 17:21:47.240842: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
113/113 [=====] - 3s 19ms/step
```



```
[57]: mse, mae = compute_metrics(series_valid, rnn_forecast)

print(f"mse: {mse:.2f}, mae: {mae:.2f} for forecast")
```

```
mse: 5.66, mae: 1.87 for forecast
```

To pass this assignment your forecast should achieve a MSE of 6 or less and a MAE of 2 or less.

- If your forecast didn't achieve this threshold try re-training your model with a different architecture (you will need to re-run both `create_uncompiled_model` and `create_model` functions) or tweaking the optimizer's parameters.
- If your forecast did achieve this threshold run the following cell to save the model in the SavedModel format which will be used for grading and after doing so, submit your assignment for grading.
- This environment includes a dummy SavedModel directory which contains a dummy model trained for one epoch. **To replace this file with your actual model you need to run the next cell before submitting for grading.**

```
[58]: # Save your model in the SavedModel format
model.save('saved_model/my_model')

# Compress the directory using tar
! tar -czvf saved_model.tar.gz saved_model/
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
lstm_cell_14_layer_call_fn,
lstm_cell_14_layer_call_and_return_conditional_losses,
lstm_cell_15_layer_call_fn,
lstm_cell_15_layer_call_and_return_conditional_losses while saving (showing 5 of
5). These functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: saved_model/my_model/assets
```

```
INFO:tensorflow:Assets written to: saved_model/my_model/assets
```

```
a saved_model
a saved_model/my_model
a saved_model/my_model/keras_metadata.pb
a saved_model/my_model/variables
a saved_model/my_model/saved_model.pb
a saved_model/my_model/assets
a saved_model/my_model/variables/variables.data-00000-of-00001
a saved_model/my_model/variables/variables.index
```

Congratulations on finishing this week's assignment!

You have successfully implemented a neural network capable of forecasting time series leveraging a combination of Tensorflow's layers such as Convolutional and LSTMs! This resulted in a forecast that surpasses all the ones you did previously.

By finishing this assignment you have finished the specialization! Give yourself a pat on the back!!!